

# Towards performance tuning mindset



Presenter Viorel Chelaru

Senior Java Developer @ METRO Systems

# About me

- Viorel Chelaru

- Mathematics & Computer Science @ UGAL
- Master in Artificial Intelligence @ PUB
- Java developer for 15 years
- Challenging projects keep me alive

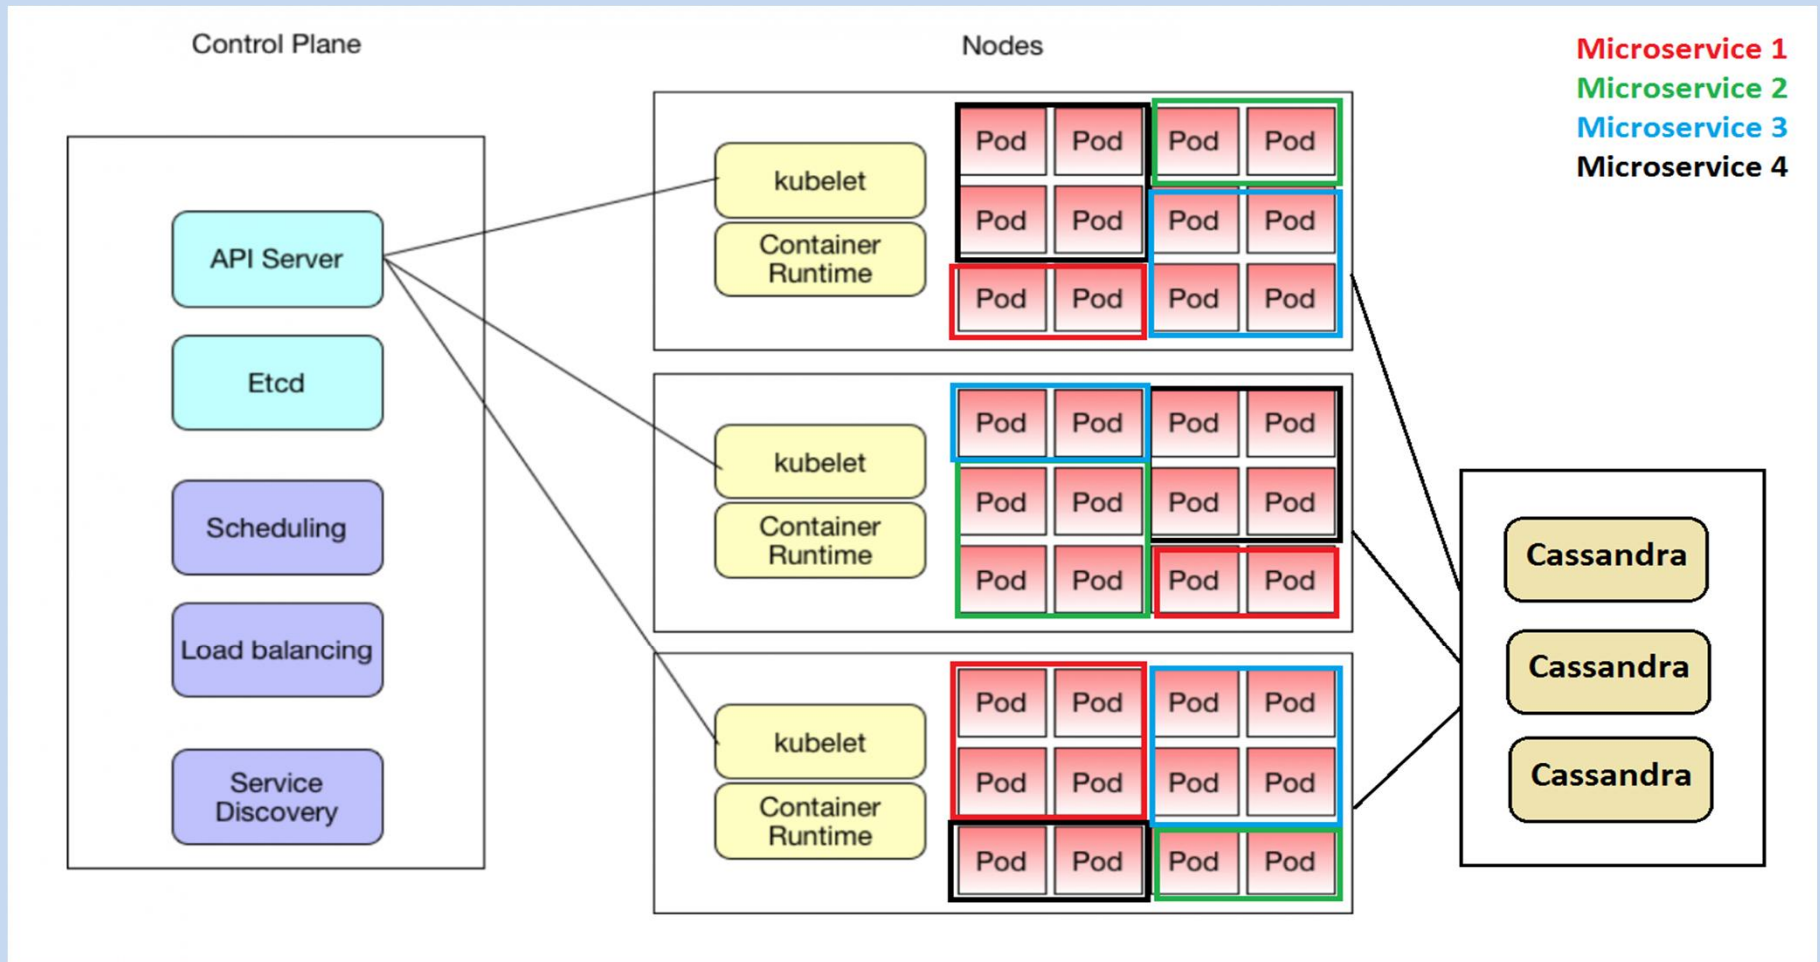
# Do we really need fast Java code ?

- I never heard a client say:
  - I am happy the 100 lines CSV upload takes 30 minutes.
- Black Friday : websites slow, until they go down
- Cashier desk in store: card payment takes 10 minutes
-

# Do we really need fast Java code ?

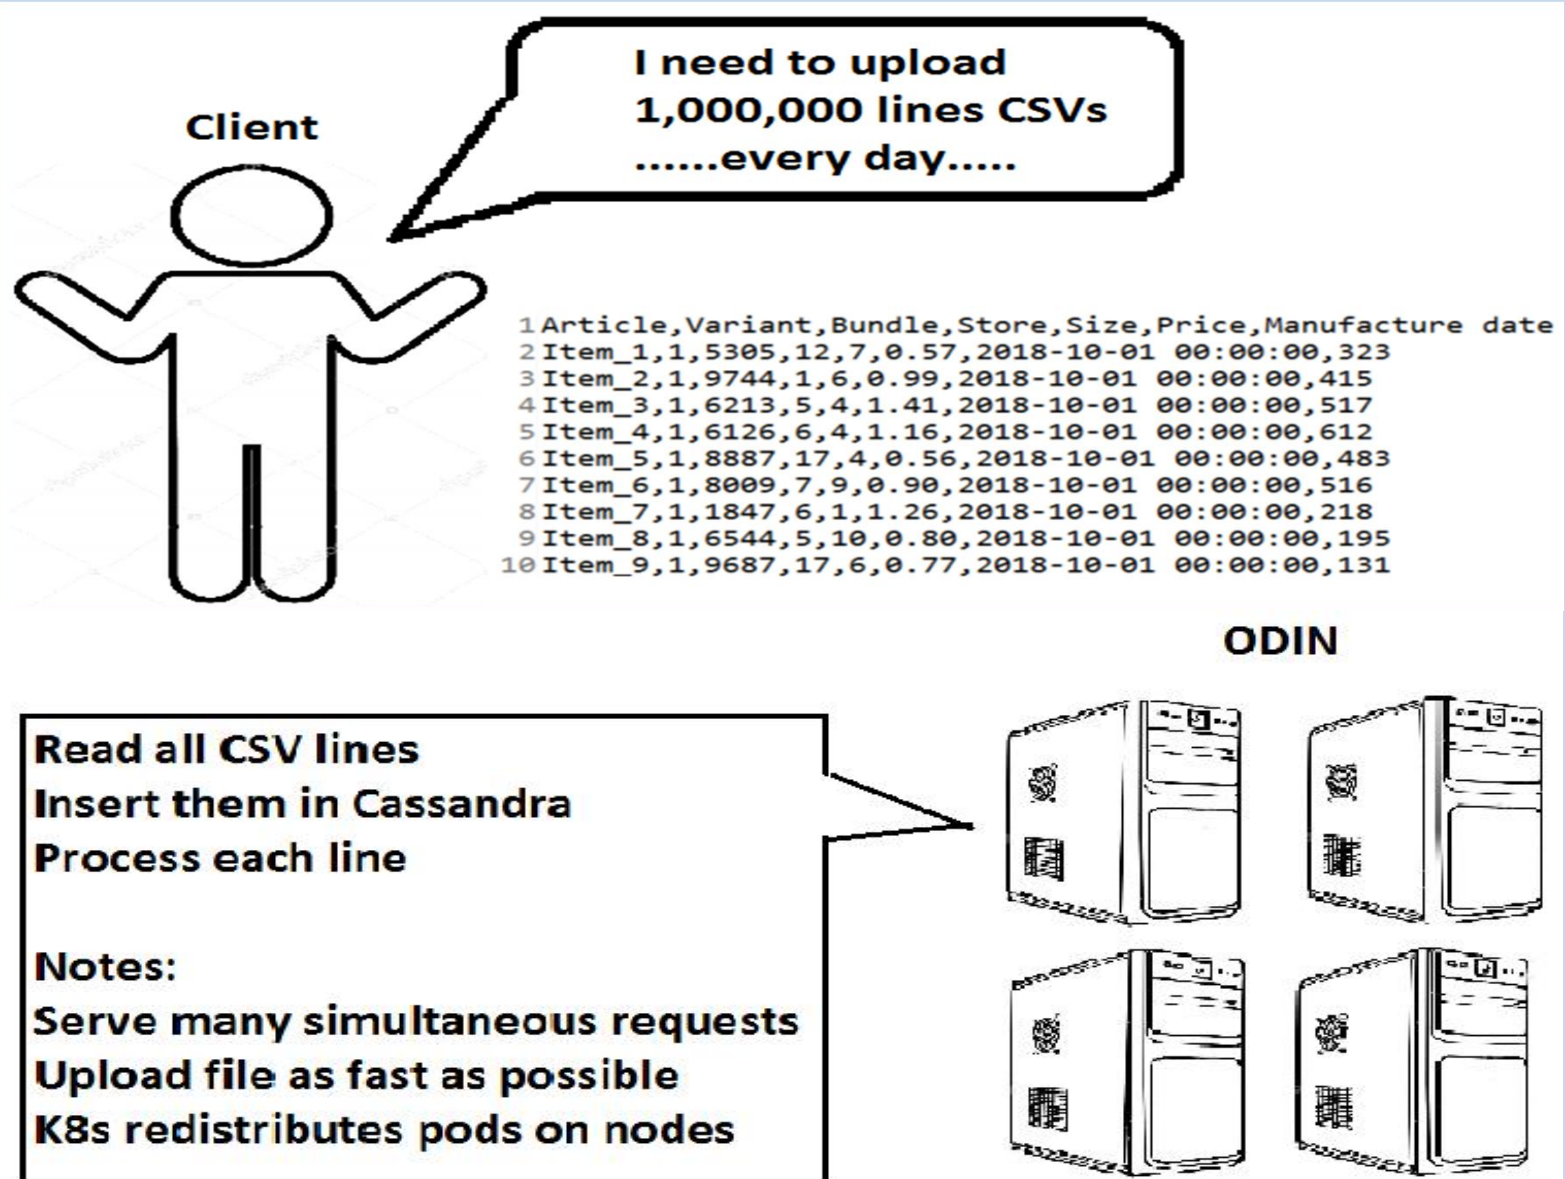
- I never heard a client say:
  - I am happy the 100 lines CSV upload takes 30 minutes.
- Black Friday : websites slow, until they go down
- Cashier desk in store: card payment takes 10 minutes
- Software performance affects our lives

# ODIN project @ METRO



- Task processing system
- REST uService in SpringBoot
- Deployed in Docker on Kubernetes
- 14 uServices
- ~200 pods on 15 nodes
- 2-80 pods per uService

# Use case : Big file upload



# It's easy to implement a file upload

```
Path path = Paths.get(DataUtil.DATA_FILEPATH);

try (Cluster cluster = DbUtil.getCluster(null);
     Session session = cluster.connect("od_inbound_dev");
     BufferedReader br = Files.newBufferedReader(path, Charset.forName("UTF-8"));) {

    PreparedStatement psInsert = session.prepare(CQL_INSERT);
    UUID fileId = UUIDs.timeBased();

    //skip header
    br.readLine();

    // read the file line by line, and insert into Cassandra
    String line = null;
    while ((line = br.readLine()) != null) {
        if (line.trim().isEmpty()) {
            continue;
        }

        String[] columns = line.split(",");
        String article = columns[0];
        int variant = Integer.parseInt(columns[1]);
        int bundle = Integer.parseInt(columns[2]);
        int store = Integer.parseInt(columns[3]);
        int size = Integer.parseInt(columns[4]);
        float price = Float.parseFloat(columns[5]);
        Date manufactureDate = sdf.parse(columns[6]);
        int valabilityDays = Integer.parseInt(columns[7]);

        session.execute(psInsert.bind(fileId, article, variant, bundle, store, size, price, manufactureDate, valabilityDays));
    }

} catch (Exception e) {
    e.printStackTrace();
    throw e;
}
```

read CSV file  
iterate all lines

parse data line

bind data in PreparedStatement  
execute PreparedStatement



# It's easy ??? to implement a file upload

```
Path path = Paths.get(DataUtil.DATA_FILEPATH);

try (Cluster cluster = DbUtil.getCluster(null);
     Session session = cluster.connect("od_inbound_dev");
     BufferedReader br = Files.newBufferedReader(path, Charset.forName("UTF-8"));) {

    PreparedStatement psInsert = session.prepare(CQL_INSERT);
    UUID fileId = UUIDs.timeBased();

    //skip header
    br.readLine();

    // read the file line by line, and insert into Cassandra
    String line = null;
    while ((line = br.readLine()) != null) {
        if (line.trim().isEmpty()) {
            continue;
        }


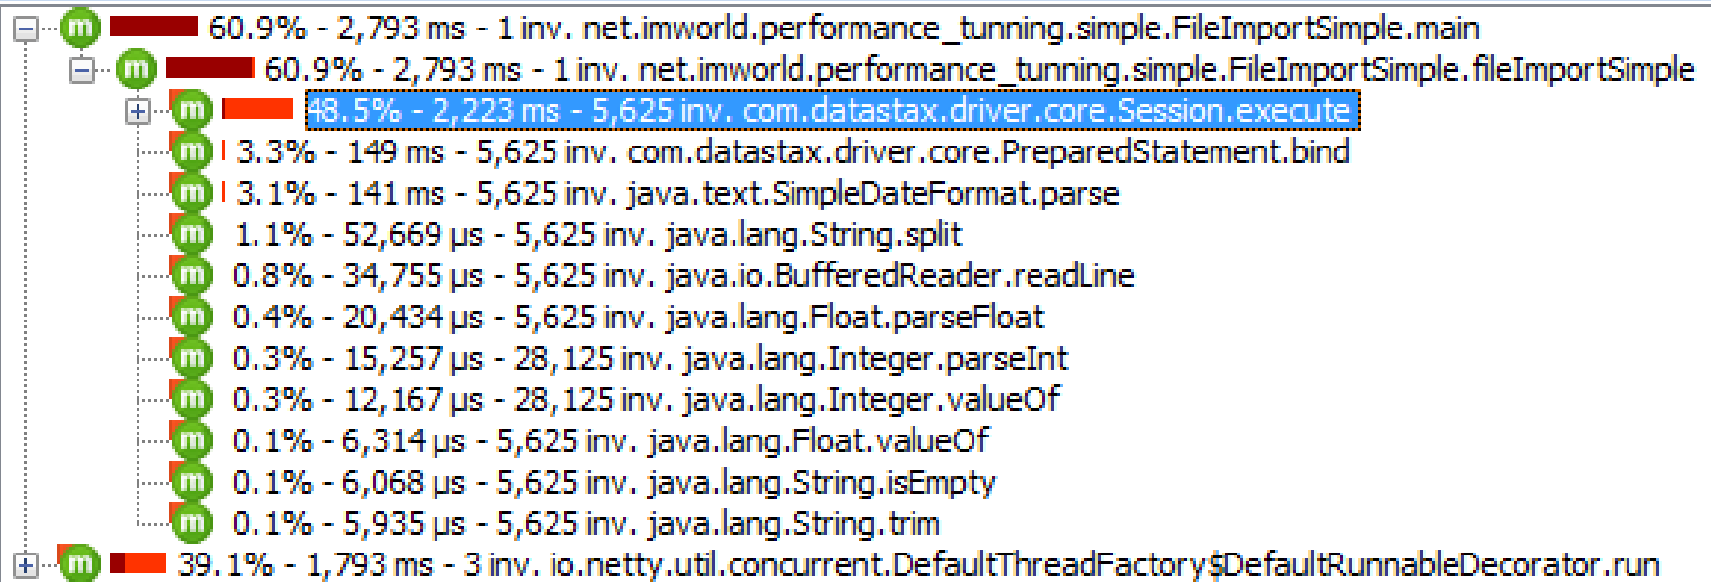
        String[] columns = line.split(",");
        String article = columns[0];
        int variant = Integer.parseInt(columns[1]);
        int bundle = Integer.parseInt(columns[2]);
        int store = Integer.parseInt(columns[3]);
        int size = Integer.parseInt(columns[4]);
        float price = Float.parseFloat(columns[5]);
        Date manufactureDate = sdf.parse(columns[6]);
        int valabilityDays = Integer.parseInt(columns[7]);

        session.execute(psInsert.bind(fileId, article, variant, bundle, store, size, price, manufactureDate, valabilityDays));
    }
} catch (Exception e) {
    e.printStackTrace();
    throw e;
}
```

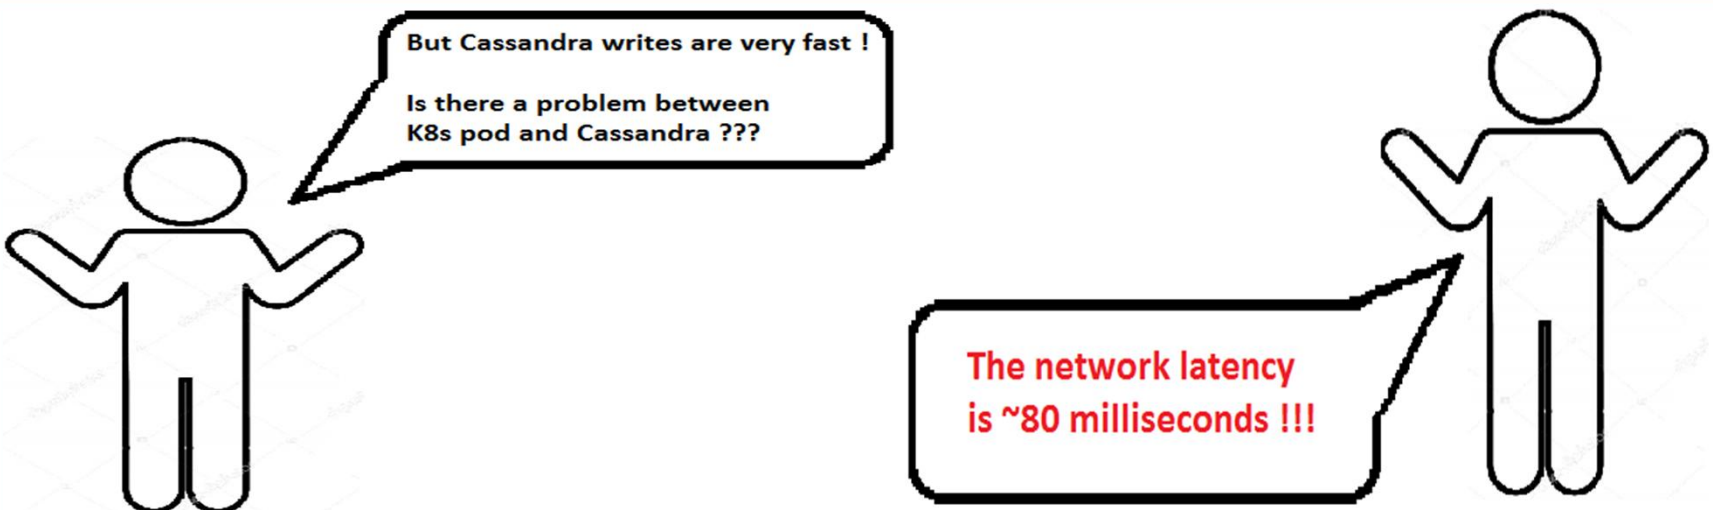
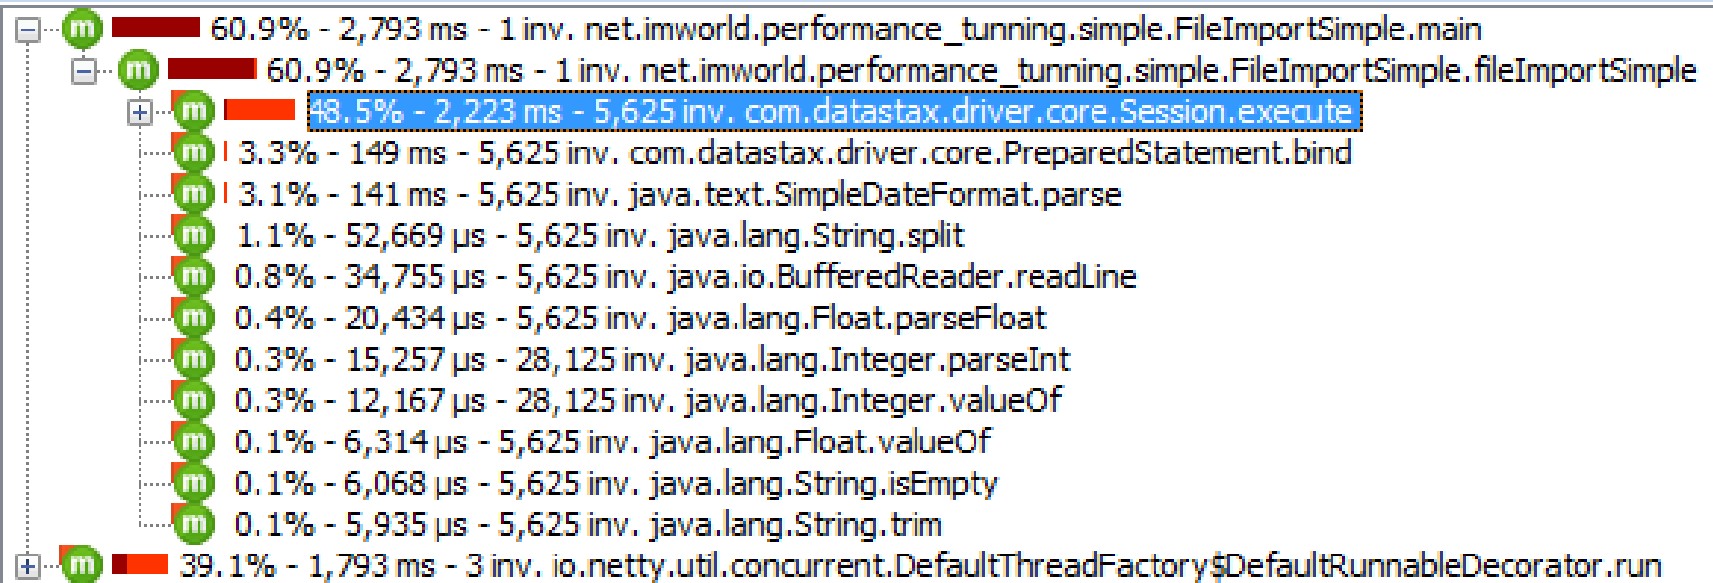
- Runs in ~22 hours
- 30Mb memory footprint



# Profiling to find the cause



# Profiling to find the cause



# Multi-threading is faster !

```
// read the file line by line, and insert into Cassandra
String line = null;
int count = 0;
List<String> linesBuffer = new LinkedList<>();
while ((line = br.readLine()) != null) {
    if (line.trim().isEmpty()) {
        continue;
    }

    count++;
    linesBuffer.add(line);

    if (count % linesPerThread == 0) {
        futures.add(processLinesBuffer(linesBuffer, fileId));
        linesBuffer.clear();
    }
}

if (!linesBuffer.isEmpty()) {
    futures.add(processLinesBuffer(linesBuffer, fileId));
}

//wait for all tasks to finish
for (Future<?> future : futures) {
    future.get();
}
```

```
private Future<?> processLinesBuffer(List<String> linesBuffer, UUID fileId) {
    SimpleDateFormat sdf = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss");
    List<String> threadLinesBuffer = new LinkedList<>();
    threadLinesBuffer.addAll(linesBuffer);
    Runnable r = new Runnable() {

        @Override
        public void run() {

            for (String line : threadLinesBuffer) {
                String[] columns = line.split(",");
                String article = columns[0];
                int variant = Integer.parseInt(columns[1]);
                int bundle = Integer.parseInt(columns[2]);
                int store = Integer.parseInt(columns[3]);
                int size = Integer.parseInt(columns[4]);
                float price = Float.parseFloat(columns[5]);
                Date manufactureDate;
                try {
                    manufactureDate = sdf.parse(columns[6]);
                } catch (ParseException e) {
                    throw new RuntimeException(e);
                }
                int valabilityDays = Integer.parseInt(columns[7]);

                session.execute(psInsert.bind(fileId, article, variant,
                    bundle, store, size, price, manufactureDate, valabilityDays));
            }
        }
    };
    return taskExecutor.submit(r);
}
```

# Multi-threading is faster !

```
// read the file line by line, and insert into Cassandra
String line = null;
int count = 0;
List<String> linesBuffer = new LinkedList<>();
while ((line = br.readLine()) != null) {
    if (line.trim().isEmpty()) {
        continue;
    }

    count++;
    linesBuffer.add(line);

    if (count % linesPerThread == 0) {
        futures.add(processLinesBuffer(linesBuffer, fileId));
        linesBuffer.clear();
    }
}

if (!linesBuffer.isEmpty()) {
    futures.add(processLinesBuffer(linesBuffer, fileId));
}

//wait for all tasks to finish
for (Future<?> future : futures) {
    future.get();
}
```

```
private Future<?> processLinesBuffer(List<String> linesBuffer, UUID fileId) {
    SimpleDateFormat sdf = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss");
    List<String> threadLinesBuffer = new LinkedList<>();
    threadLinesBuffer.addAll(linesBuffer);
    Runnable r = new Runnable() {

        @Override
        public void run() {

            for (String line : threadLinesBuffer) {
                String[] columns = line.split(",");
                String article = columns[0];
                int variant = Integer.parseInt(columns[1]);
                int bundle = Integer.parseInt(columns[2]);
                int store = Integer.parseInt(columns[3]);
                int size = Integer.parseInt(columns[4]);
                float price = Float.parseFloat(columns[5]);
                Date manufactureDate;
                try {
                    manufactureDate = sdf.parse(columns[6]);
                } catch (ParseException e) {
                    throw new RuntimeException(e);
                }
                int valabilityDays = Integer.parseInt(columns[7]);


                session.execute(psInsert.bind(fileId, article, variant,
                    bundle, store, size, price, manufactureDate, valabilityDays));
            }
        }
    };
    return taskExecutor.submit(r);
}
```

2 parameters to adjust


- #lines per thread
- #threads in the pool

# Multi-threading results

Running time in minutes



Memory footprint in Mb



Multi-threading solution:


- read 1,000,000 lines with a pool of 80 threads, each thread inserting 10,000 lines
- Running time 23 mins
- Memory footprint 400 Mb

Using Cassandra SimpleStatement (CQL concat), not binding to a PreparedStatement


- Running time 18 mins

# Multi-threading results

Running time in minutes



Memory footprint in Mb



Multi-threading solution:

- read 1,000,000 lines with a pool of 80 threads, each thread inserting 10,000 lines
- Running time 23 mins
- Memory footprint 400 Mb

Using Cassandra SimpleStatement (CQL concat), not binding to a PreparedStatement

- Running time 18 mins

Still NOT HAPPY ! 18 minutes is too much !  
1 pod has 450Mb required memory !



# Inserting with Cassandra batches

```
// read the file line by line, and insert into Cassandra
String line = null;
List<String> linesBuffer = new ArrayList<>();
while ((line = br.readLine()) != null) {
    if (line.trim().isEmpty()) {
        continue;
    }

    linesBuffer.add(line);

    if (linesBuffer.size() == linesBufferSize) {
        insertLinesBuffer(fileId, linesBuffer, batchSize);
        linesBuffer.clear();
    }

    if (!linesBuffer.isEmpty()) {
        insertLinesBuffer(fileId, linesBuffer, batchSize);
    }
}
```

```
private void insertLinesBuffer(UUID fileId, List<String> linesBuffer, int batchSize)
    throws ParseException {
    //logger.info("insert buffer of "+linesBuffer.size()+" lines");

    BatchStatement batch = new BatchStatement();
    for (String line : linesBuffer) {
        String[] columns = line.split(",");
        String article = columns[0];

        int variant = Integer.parseInt(columns[1]);
        int bundle = Integer.parseInt(columns[2]);
        int store = Integer.parseInt(columns[3]);
        int size = Integer.parseInt(columns[4]);
        float price = Float.parseFloat(columns[5]);
        Date manufactureDate = sdf.parse(columns[6]);
        int valabilityDays = Integer.parseInt(columns[7]);

        batch.add(psInsert.bind(fileId, article, variant, bundle, store, size,
            price, manufactureDate, valabilityDays));

        if (batch.size() == batchSize) {
            session.execute(batch);
            batch.clear();
        }
    }

    if (batch.size() > 0) {
        session.execute(batch);
        batch.clear();
    }
}
```


Benefit:

send multiple statements in 1 network call




# Inserting with Cassandra batches

Running time in minutes



Memory footprint in Mb



Cassandra Batches solution:


- read 1,000,000 lines in buffers of 1000 lines, insert buffers in batches of 400 inserts
- Running time 5 mins
- Memory footprint 50 Mb

Using Cassandra SimpleStatement (CQL concat), not binding to a PreparedStatement


- Running time 11 mins

# Inserting with Cassandra batches

Running time in minutes



Memory footprint in Mb



Cassandra Batches solution:

- read 1,000,000 lines in buffers of 1000 lines, insert buffers in batches of 400 inserts
- Running time 5 mins
- Memory footprint 50 Mb

Using Cassandra SimpleStatement (CQL concat), not binding to a PreparedStatement

- Running time 11 mins

We're getting somewhere !  
5 minutes is pretty good, but can I do better ?

# Does Garbage Collection count ?



Cassandra Batches solution:

- read 1,000,000 lines in buffers of 1000 lines
- insert buffers in batches of 400 inserts
- Use G1 garbage collector
- Running time 4 mins
- Memory footprint 50 – 80 Mb

# Does Garbage Collection count ?



Cassandra Batches solution:

- read 1,000,000 lines in buffers of 1000 lines
- insert buffers in batches of 400 inserts
- Use G1 garbage collector
- Running time 4 mins
- Memory footprint 50 – 80 Mb

Wowwww, I just squized 1 MINUTE !  
4 minutes is niceeee....  
I can see the end of the tunnel !

# 3000 batches ? Should be 2500 batches !

Until now

1,000,000 lines read in buffers of 1000 lines

Insert each buffer in 400 inserts batches

1000 lines

1 buffer => 3 batches (400,400,200)

1,000,000 lines

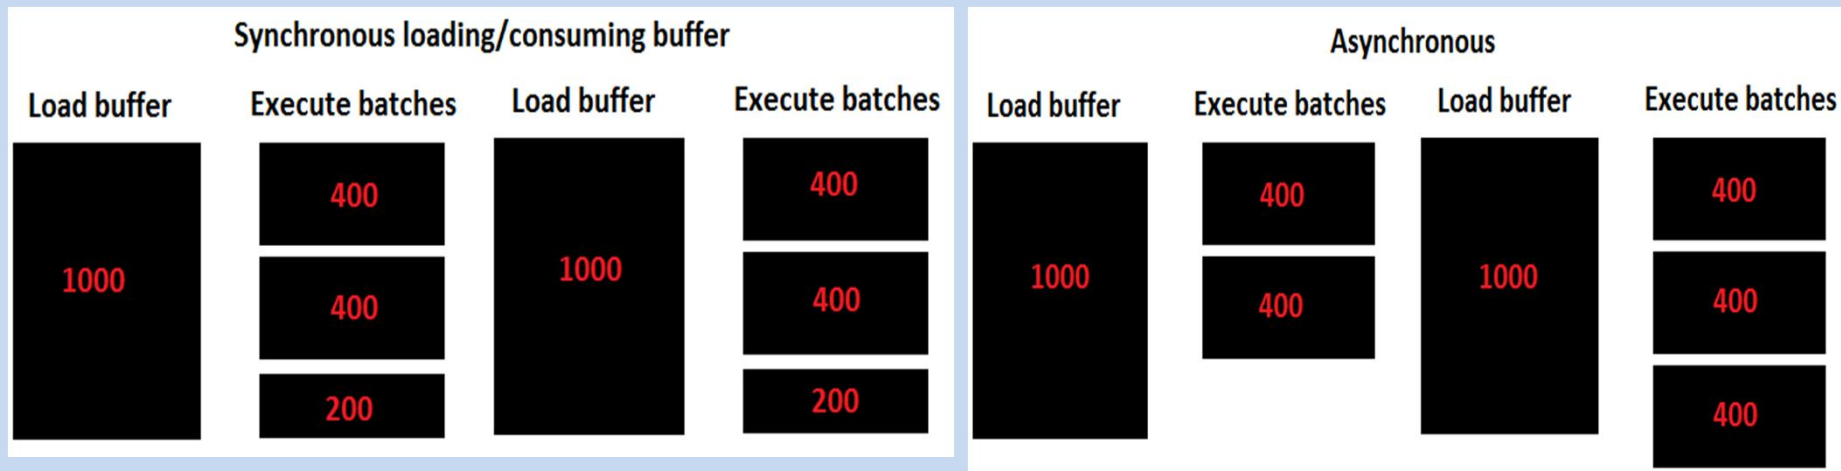
1000 buffers \* 3 batches / buffer = 3000 batches

But

1,000,000 lines / 400 inserts = **2500 batches !!!!**

Does it worth optimizing ?

500 batches \* 80ms latency / batch = 40 seconds



# Final solution

- Use Cassandra batches (single partition)
- read 1,000,000 lines in buffers of 1000 lines
- insert buffers in batches of 400 inserts
- Decouple loading buffer and executing insert batches
- Use G1 garbage collector

Running time 3 mins


Memory footprint 50 – 80 Mb

From 22 hours to 3 minutes ??? What a journey !

We improve 440 times !!

I can go like this “forever” J

# Performance & Scalability





# I don't have to be a genius to fine tune performance

I am restless, ambitious, skeptical, always unsatisfied !

I always keep performance in mind when I code.

I HAVE time for performance tuning, if I prioritize my work.

I use my analytical thinking, not my typing skills.

THINK !

IT AIN'T ILLEGAL .... YET

Eddie Griffin, The comedian

# Thank You !

- Questions ?
- Code is on GitHub, look me up